

Essential C Security 102

Offensive Computer Security
Florida State University
Spring 2014

Outline of Talk

- Continuation of Heap / Dynamic Memory discussion
- Integer Security
- Formatted Output
- Concurrency and Race Conditions

Tool we will be using

<http://gcc.godbolt.org/>

A project that visualizes C/C++ to Assembly for you. *(use g++ compiler, intel syntax, and no options)*

Quite useful for learning this stuff
(also interesting: <https://github.com/ynh/cpp-to-assembly>)

Memory Allocator

The memory manager on most systems runs as part of the process

- linker adds in code to do this
 - usually provided to linker via OS
 - OS's have default memory managers
 - compilers can override or provide alternatives
- Can be statically linked in or dynamically

Done! unlinked!

```
// This moves a chunk from  
// the free list, to be used
```

```
#define unlink(P, BK, FD) {  
    FD = P->fd;  
    BK = P->bk;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```

BK

Free list

FD

| | | | | |
|---|--|--|--|--------------|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |
| unused | | | | |
| Size of this chunk | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |
| unused | | | | |
| Size of this chunk | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |

Allocated chunks

| | | | | |
|----------------------------------|--|--|--|---|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |

Memory Allocator

In general requires:

- A maintained list of free, available memory
- algorithm to allocate a contiguous chunk of n bytes
 - Best fit method
 - chunk of size $m \geq n$ such that m is smallest available
 - First fit method
- algorithm to deallocate said chunks (free)
 - return chunk to list, consolidate adjacent used ones.

Memory Allocator

Common optimizations:

- Chunk boundary tags

- [tag][-----chunk -----][tag]

- tag contains metadata:

- size of chunk
 - next chunk
 - previous chunk (like a linked list sometimes)

Doug Lea's dmalloc allocator

Basis of Linux mem allocators

- Free chunks are arranged in a doubly-linked circular lists (**bins**)
- Each chunk (used and free) has:
 - next chunk and previous chunk pointers
 - size of previous chunk (if free) / last 4 bytes of the previous used chunk (if not free)
 - *Last 4 bytes is a mystery to me, don't ask me*
 - flag for if previous chunk is used / free

Doug Lea's dlmalloc allocator

Allocated chunk

| | | | | |
|----------------------------------|--|--|--|---|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |
| data | | | | |
| | | | | |
| Last 4 bytes of user data | | | | |
| Size of next chunk | | | | 1 |

chunk
boundaries

Chunk
1

Chunk
2

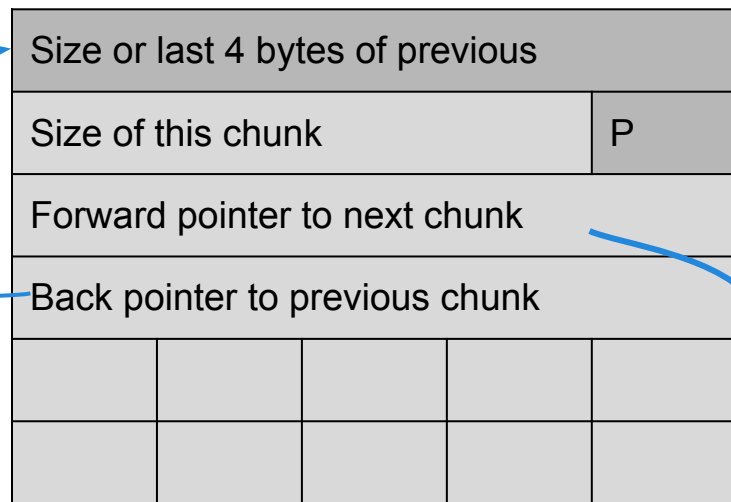
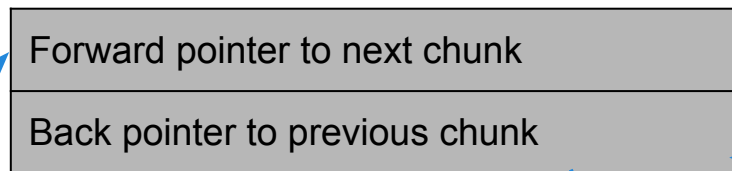
Free chunk

| | | | | |
|----------------------------------|--|--|--|---|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |
| | | | | |
| | | | | |
| Size of this chunk | | | | |
| Size of next chunk | | | | 0 |

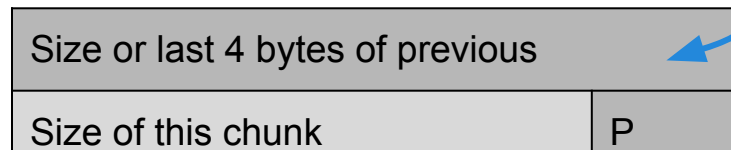
Doug Lea's dmalloc allocator

Each list of free bin has a head: **Free chunk**

- doubly linked list
 - forward pointer
 - backwards pointer



■ ■ ■



How unlinking works

```
// This moves a chunk from  
// the free list, to be used
```

```
#define unlink(P, BK, FD) {  
    FD = P->fd;  
    BK = P->bk;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```

```
//This is from [1]p 184
```

Free list

| | |
|----------------------------------|---|
| Size or last 4 bytes of previous | |
| Size of this chunk | P |
| Forward pointer to next chunk | |
| Back pointer to previous chunk | |
| unused | |
| Size of this chunk | |
| Size or last 4 bytes of previous | |
| Size of this chunk | P |
| Forward pointer to next chunk | |
| Back pointer to previous chunk | |
| unused | |
| Size of this chunk | |
| Size or last 4 bytes of previous | |
| Size of this chunk | P |
| Forward pointer to next chunk | |
| Back pointer to previous chunk | |

Allocated chunks

| | | | | |
|----------------------------------|--|--|--|---|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |

Say this is P

```
// This moves a chunk from  
// the free list, to be used
```

```
#define unlink(P, BK, FD) {  
    FD = P->fd;  
    BK = P->bk;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```

Free list

| | |
|----------------------------------|---|
| Size or last 4 bytes of previous | |
| Size of this chunk | P |
| Forward pointer to next chunk | |
| Back pointer to previous chunk | |
| unused | |
| Size of this chunk | |
| Size or last 4 bytes of previous | |
| Size of this chunk | P |
| Forward pointer to next chunk | |
| Back pointer to previous chunk | |
| unused | |
| Size of this chunk | |
| Size or last 4 bytes of previous | |
| Size of this chunk | P |
| Forward pointer to next chunk | |
| Back pointer to previous chunk | |

Allocated chunks

| | | | | |
|----------------------------------|--|--|--|---|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |

1) $FD = P \rightarrow fd;$

// This moves a chunk from
// the free list, to be used

```
#define unlink(P, BK, FD) {  
    FD = P->fd;  
    BK = P->bk;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```

Setting this
to FD

Free list

| | |
|----------------------------------|---|
| Size or last 4 bytes of previous | |
| Size of this chunk | P |
| Forward pointer to next chunk | |
| Back pointer to previous chunk | |
| unused | |
| Size of this chunk | |
| Size or last 4 bytes of previous | |
| Size of this chunk | P |
| Forward pointer to next chunk | |
| Back pointer to previous chunk | |
| unused | |
| Size of this chunk | |
| Size or last 4 bytes of previous | |
| Size of this chunk | P |
| Forward pointer to next chunk | |
| Back pointer to previous chunk | |

Allocated chunks

| | | | | |
|----------------------------------|--|--|--|---|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |

2) BK = P->bk;

// This moves a chunk from
// the free list, to be used

```
#define unlink(P, BK, FD) {
    FD = P->fd;
    BK = P->bk;
    FD->bk = BK;
    BK->fd = FD;
}
```

Setting
this to BK

Free list

| | | | | |
|----------------------------------|--|--|--|---|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |
| unused | | | | |
| Size of this chunk | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |
| unused | | | | |
| Size of this chunk | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |

Allocated chunks

| | | | | |
|----------------------------------|--|--|--|---|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |

FD

3) $FD \rightarrow bk = BK;$

// This moves a chunk from
// the free list, to be used

```
#define unlink(P, BK, FD) {  
    FD = P->fd;  
    BK = P->bk;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```

BK

Free list

Allocated chunks

| | |
|----------------------------------|---|
| Size or last 4 bytes of previous | |
| Size of this chunk | P |
| Forward pointer to next chunk | |
| Back pointer to previous chunk | |
| unused | |
| Size of this chunk | |
| Size or last 4 bytes of previous | |
| Size of this chunk | P |
| Forward pointer to next chunk | |
| Back pointer to previous chunk | |
| unused | |
| Size of this chunk | |
| Size or last 4 bytes of previous | |
| Size of this chunk | P |
| Forward pointer to next chunk | |
| Back pointer to previous chunk | |

| | | | | |
|----------------------------------|--|--|--|---|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |

FD

4) BK->fd = FD;

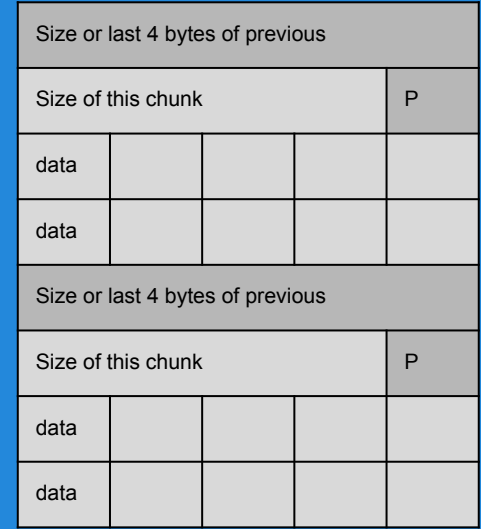
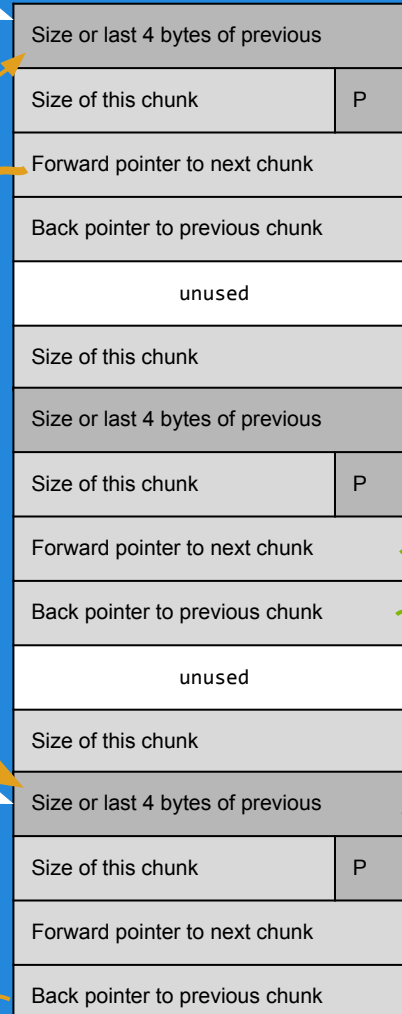
// This moves a chunk from
// the free list, to be used

```
#define unlink(P, BK, FD) {  
    FD = P->fd;  
    BK = P->bk;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```

BK → **Free list**

FD →

Allocated chunks



Chunk is now Allocated

Things to note:

- Two pointers are changed
 - BK->fd
 - FD->bk
 - Keep this in mind
- This trusts the data in the system to work right
 - double malloc doesn't mess this up
 - not a bug

free() is the reverse of this process

- involves changing pointers
 - double free messes this up!

BK

Free list

| | | | | |
|---|--|--|--|---|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |
| unused | | | | |
| Size of this chunk | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |
| unused | | | | |
| Size of this chunk | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |

FD

Allocated chunks

| | | | | |
|----------------------------------|--|--|--|---|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |

Chunk is now Allocated

`free()` is the reverse of this process

- involves changing pointers
 - double free messes this up!
- Majority of buffer overflows since 2000 have been on the heap [1]
 - b/c devs don't understand it well

BK

Free list

| | | | | |
|---|--|--|--|--------------|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |
| unused | | | | |
| Size of this chunk | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |
| unused | | | | |
| Size of this chunk | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |

FD

Allocated chunks

| | | | | |
|----------------------------------|--|--|--|---|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk | | | | P |
| data | | | | |
| data | | | | |

Exploring Heap Vulnerabilities

For these examples we'll use this guy as our friendly guide

- likes `free()`[dom]
- likes heaps [of british skulls]



How free works (from [1] p186)

```
#include <stdlib.h>
#include <string.h>
int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```

why 672 not
666?

| | | | | |
|----------------------------------|--|--|--|-----|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 672 | | | | P=1 |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |

How free works (from [1] p186)

```
#include <stdlib.h>
#include <string.h>
int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```

| | | | | |
|----------------------------------|--|--|--|-----|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 672 | | | | P=1 |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |

How free works (from [1] p186)

```
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```

```
//pseudo code for free()
define free() {
    if (next not in use)
        consolidate with next;
        //(merges with existing chunk
        on free list)
    else
        link chunk to free list;
}
```

| | | | | |
|----------------------------------|--|--|--|-----|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 672 | | | | P=1 |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |

How free works (from [1] p186)

```
#include <stdlib.h>
#include <string.h>
int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```

Checks to see if next chunk is also free

- checks P (PREV_IN_USE) flag on next, next chunk
 - it finds this via the size metadata in the current chunk and next chunk

In this case it is in use

So first is just freed up and linked to the free list

The P flag on the next bin (second) is then set to 0

| | | | | |
|----------------------------------|-----|--|--|--|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 672 | P=1 | | | |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | P=0 | | | |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | P=1 | | | |
| data | | | | |
| data | | | | |

How free works (from [1] p186)

```
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```

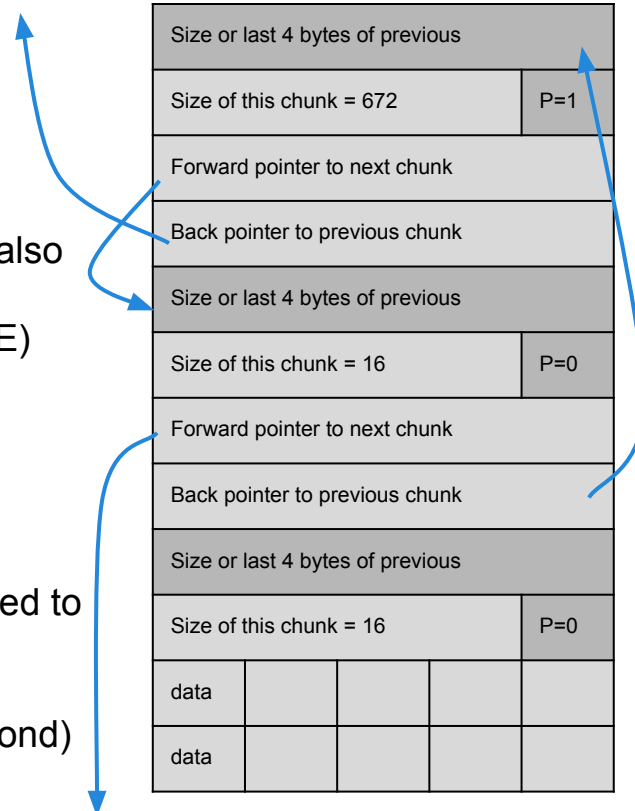
Checks to see if next chunk is also free

- checks P (PREV_IN_USE) flag on next, next chunk (not shown)

In this case it is in use

So first is just freed up and linked to the free list

The P flag on the next bin (second) is then set to 0



| | | | | |
|----------------------------------|--|--|--|-----|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 672 | | | | P=1 |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=0 |
| Forward pointer to next chunk | | | | |
| Back pointer to previous chunk | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=0 |
| data | | | | |
| data | | | | |

How free works (from [1] p186)

```
#include <stdlib.h>
#include <string.h>


int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```

and so on

Note that consolidation may happen, and this is not shown

- **consolidation calls that unlink macro we discussed earlier**




How free works (from [1] p186)

```
#include <stdlib.h>
#include <string.h>
int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```

What to note:

- Pointers changed
 - in the chunk freed
 - and in OTHER chunks!
 - relies on meta data being correct
 - lets explore how this can be subverted maliciously
 - (arbitrary memory write vuln)



How free works (from [1] p186)

```
#include <stdlib.h>
#include <string.h>
int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```

This metadata
is the target

but we can only hit the
2nd one here

| | | | | |
|----------------------------------|--|--|--|-----|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 672 | | | | P=1 |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=0 |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |

Heap Buffer Overflow (from [1] p186)

```
#include <stdlib.h>
#include <string.h>
int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```



| | | | | |
|----------------------------------|--|--|--|-----|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 672 | | | | P=1 |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |

Heap Buffer Overflow (from [1] p186)

```
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```

[illegible]

| | |
|--|-----|
| Size or last 4 bytes of previous | |
| Size of this chunk = 672 | P=1 |
|  | |

Heap Buffer Overflow (from [1] p186)

```
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```



Will cause free
(second)
to segfault

[illegible]

| | |
|--|-----|
| Size or last 4 bytes of previous | |
| Size of this chunk = 672 | P=1 |
|  | |

Heap Buffer Overflow (from [1] p186)

```
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```



will alter the
behavior of
free()

“FREEEEEEEEEEEDOOM
MMMMMMMMMMMMMMMM
MMMMMMMM...<dummy even
integer(to have P=0)><new
size (-4)><a fd pointer><a bk
pointer>”

| | | | | |
|--|--|--|--|-----|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 672 | | | | P=1 |
| FREEEEEEDOOOMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMM MMMMMMMMMMMM | | | | |
| dummy size field | | | | P=0 |
| size of chunk = -4 | | | | P=0 |
| Malicious fd pointer | | | | |
| Malicious bk pointer | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |

Heap Buffer Overflow (from [1] p186)

```
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```



Size field in second chunk
overwritten with a negative
number

- when `free()` attempts to find the third chunk it will go here:

| | | | | |
|--|--|--|--|-----|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 672 | | | | P=1 |
| FREEEEEEDOOOMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMM | | | | |
| dummy size field | | | | F=0 |
| size of chunk = -4 | | | | P=0 |
| Malicious fd pointer | | | | |
| Malicious bk pointer | | | | |
| | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |

Heap Buffer Overflow (from [1] p186)

```
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```



Size field in second chunk
overwritten with a negative
number

- when free() attempts to find the third chunk it will go here:
 - it sees the 2nd chunk is listed as free
 - unlink time

| | | | | |
|--|--|--|--|-----|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 672 | | | | P=1 |
| FREEEEEEDOOOMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMM | | | | |
| dummy size field | | | | F=0 |
| size of chunk = -4 | | | | P=0 |
| Malicious fd pointer | | | | |
| Malicious bk pointer | | | | |
| | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |

Heap Buffer Overflow (from [1] p186)

```
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[])
{
    char *first, *second, *third;
    first = malloc(666);
    second = malloc(12);
    third = malloc(12);
    strcpy(first, argv[1]);

    free(first);
    free(second);
    free(third);
}
```



```
#define unlink(P, BK, FD) {
    FD = P->fd;
    BK = P->bk;
    FD->bk = BK;
    BK->fd = FD;
}
```

need not point to
the heap or to
the free list!

| | | | | |
|--|--|--|--|-----|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 672 | | | | P=1 |
| FREEEEEEDOOOMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMM MMMMMMMMMM | | | | |
| dummy size field | | | | P=0 |
| size of chunk = -4 | | | | P=0 |
| Malicious fd pointer | | | | |
| Malicious bk pointer | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |

Heap Buffer Overflow (from [1] p186)

When this
command
runs:

- writes attacker supplied data to an attacker supplied address

- to (fd + 12)
 - why?



```
#define unlink(P, BK, FD) {
    FD = P->fd;
    BK = P->bk;
    FD->bk = BK;
    BK->fd = FD;
}
```

The destination of the arbitrary write

The value which to write

| | | | | |
|---|--|--|--|-----|
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 672 | | | | P=1 |
| FREEEEEEDOOOMMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMM MMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMM | | | | |
| dummy size field | | | | P=0 |
| size of chunk = -4 | | | | P=0 |
| Malicious fd pointer | | | | |
| Malicious bk pointer | | | | |
| Size or last 4 bytes of previous | | | | |
| Size of this chunk = 16 | | | | P=1 |
| data | | | | |
| data | | | | |

Double free() bug (kinda) does this



Take this guy `free()` him



`free()` him again and it produces
some messed up zombie state
of the former heap

Double free() Vulnerability

- Another exploitable bug
- Conditions to be vulnerable:
 - chunk to be free()'d must be isolated (no free adjacent chunks, they must be in use).
 - the free list bin in which the chunk is going must be empty (all those size-chunks must be in use)

Double free() Vulnerability

- much more complicated than the last bug
 - sadly we don't have time for it
 - See *“Secure Coding in C and C++”* by Robert Seacord for a great discussion
- affects dlmalloc and old versions of RtlHeap
 - most modern allocator alternatives do safe unlinking
 - prevents most double frees

Use after free() Vulnerability

- involves using a pointer to a heap chunk that has been freed
 - when used as a function pointer == vulnerability
 - To exploit: need to overwrite that free'd portion of memory with malicious substitute

Integer Security

- Signed vs Unsigned
- Integer Truncation
- Overflow
- Underflow
- Nuances
- Conversion / Promotion
- Casting

Truncation Example 1

```
#include <stdio.h>
void foo()
{
    int i = -1;
    short x;
    x = i;
}
```

Lets see how this compiles and exactly what happens

| | | | | | |
|----------------|-----|----|----|---|---|
| 31 | 8 | 15 | 8 | 7 | 0 |
| Alternate name | AX | | | | |
| | AH | | AL | | |
| | EAX | | | | |
| Alternate name | BX | | | | |
| | BH | | BL | | |
| | EBX | | | | |
| Alternate name | CX | | | | |
| | CH | | CL | | |
| | ECX | | | | |
| Alternate name | DX | | | | |
| | DH | | DL | | |
| | EDX | | | | |
| Alternate name | BP | | | | |
| | EBP | | | | |
| Alternate name | SI | | | | |
| | ESI | | | | |
| Alternate name | DI | | | | |
| | EDI | | | | |
| Alternate name | SP | | | | |
| | ESP | | | | |

Truncation Example 1

Code editor

```
1 #include <stdio.h>
2 void foo()
3 {
4     int i = -1;
5     short x;
6     x = i;
7 }
8
```

Assembly output

```
1 foo():
2     push    rbp
3     mov     rbp, rsp
4     mov     DWORD PTR [rbp-4], -1
5     mov     eax, DWORD PTR [rbp-4]
6     mov     WORD PTR [rbp-6], ax
7     pop     rbp
8     ret
9
```

| 31 | 8 | 15 | 8 | 7 | 0 |
|----------------|-----|----|----|---|---|
| Alternate name | AX | | | | |
| | AH | | AL | | |
| | EAX | | | | |
| Alternate name | BX | | | | |
| | BH | | BL | | |
| | EBX | | | | |
| Alternate name | CX | | | | |
| | CH | | CL | | |
| | ECX | | | | |
| Alternate name | DX | | | | |
| | DH | | DL | | |
| | EDX | | | | |
| Alternate name | BP | | | | |
| | EBP | | | | |
| Alternate name | SI | | | | |
| | ESI | | | | |
| Alternate name | DI | | | | |
| | EDI | | | | |
| Alternate name | SP | | | | |
| | ESP | | | | |

x86_64 vs x86_32

| 64-bit register | Lower 32 bits | Lower 16 bits | Lower 8 bits |
|-----------------|---------------|---------------|--------------|
| rax | eax | ax | al |
| rbx | ebx | bx | bl |
| rcx | ecx | cx | cl |
| rdx | edx | dx | dl |
| rsi | esi | si | sil |
| rdi | edi | di | dil |
| rbp | ebp | bp | bpl |
| rsp | esp | sp | spl |
| r8 | r8d | r8w | r8b |
| r9 | r9d | r9w | r9b |
| r10 | r10d | r10w | r10b |
| r11 | r11d | r11w | r11b |
| r12 | r12d | r12w | r12b |
| r13 | r13d | r13w | r13b |
| r14 | r14d | r14w | r14b |
| r15 | r15d | r15w | r15b |

| | | | | | |
|----------------|----|----|----|---|---|
| 31 | 8 | 15 | 8 | 7 | 0 |
| Alternate name | AX | | | | |
| | AH | | AL | | |
| EAX | | | | | |
| Alternate name | BX | | | | |
| | BH | | BL | | |
| EBX | | | | | |
| Alternate name | CX | | | | |
| | CH | | CL | | |
| ECX | | | | | |
| Alternate name | DX | | | | |
| | DH | | DL | | |
| EDX | | | | | |
| Alternate name | BP | | | | |
| EBP | | | | | |
| Alternate name | SI | | | | |
| ESI | | | | | |
| Alternate name | DI | | | | |
| EDI | | | | | |
| Alternate name | SP | | | | |
| ESP | | | | | |

Integer Truncation continued

Do not code your applications with the native C/C++ data types that change size on a 64-bit operating system

- use type definitions or macros that explicitly call out the size and type of data contained in a variable

The 64-bit return value from `sizeof` in the following statement is truncated to 32-bits when assigned to `bufferSize`.

```
int bufferSize = (int) sizeof (something);
```

The solution is to cast the return value using `size_t` and assign it to `bufferSize` declared as `size_t` as shown below:

```
size_t bufferSize = (size_t) sizeof (something);
```

Safe type definitions/functions

- **ptrdiff_t**: A signed integer type that results from subtracting two pointers.
- **size_t**: An unsigned integer and the result of the sizeof operator. This is used when passing parameters to functions such as malloc (3), and returned from several functions such as fread (2).
- **int32_t, uint32_t etc.:** Define integer types of a predefined width.
- **intptr_t and uintptr_t**: Define integer types to which any valid pointer to void can be converted.

Platform Matters (intro)

- In many programming environments for C and C-derived languages on 64-bit machines, "int" variables are still 32 bits wide
 - but long integers and pointers are 64 bits wide.
 - This is described as the LP64 data model
- Alternative models:
 - ILP64 (all 3 types are 64 bits wide)
 - SILP64 (even shorts are 64 bits wide)
 - LLP64 (compatibility mode, everything is 32 bit)

Platform Matters

The difference among the three 64-bit models (LP64, LLP64, and ILP64) lies in the non-pointer data types.

ILP32 = Microsoft Windows & Most Unix and Unix-like systems @ 32bit

LP64 = Most [Unix](#) and [Unix-like](#) systems, e.g. [Solaris](#), [Linux](#), [BSD](#), and [OS X](#); [z/OS](#)

LLP64 = [Microsoft Windows](#) (x86-64 and IA-64)

ILP64 = [HAL Computer Systems](#) port of Solaris to [SPARC64](#)

Table 1. 32-bit and 64-bit data models

| | ILP32 | LP64 | LLP64 | ILP64 |
|-----------|-------|------|-------|-------|
| char | 8 | 8 | 8 | 8 |
| short | 16 | 16 | 16 | 16 |
| int | 32 | 32 | 32 | 64 |
| long | 32 | 64 | 32 | 64 |
| long long | 64 | 64 | 64 | 64 |
| pointer | 32 | 64 | 64 | 64 |

A side note on Exploit Dev

The size of a struct may change from platform to platform!

```
struct test {  
    int i1;  
    double d;  
    int i2;  
    long l;  
}
```

Why does this matter?

Can you find the bug? (60 secs)

```
__inline__ unsigned long long int rdtsc()
{
#ifdef __x86_64__
    unsigned int a, d;
    __asm__ volatile ("rdtsc" : "=a" (a), "=d" (d));
    return (unsigned long)a | ((unsigned long)d << 32);
#elif defined(__i386__)
    unsigned long long int x;
    __asm__ volatile ("rdtsc" : "=A" (x));
    return x;
#else
#define NO_CYCLE_COUNTER
    return 0;
#endif
}
```

Table 1. 32-bit and 64-bit data models

| | ILP32 | LP64 | LLP64 | ILP64 |
|-----------|-------|------|-------|-------|
| char | 8 | 8 | 8 | 8 |
| short | 16 | 16 | 16 | 16 |
| int | 32 | 32 | 32 | 64 |
| long | 32 | 64 | 32 | 64 |
| long long | 64 | 64 | 64 | 64 |
| pointer | 32 | 64 | 64 | 64 |

Another Example

- On a 32-bit system, `int` and `long` are of the same size.
 - Due to this, some developers use them interchangeably. This can cause pointers to be assigned to `int` and vice-versa.
 - But on a 64-bit system, assigning a pointer to an `int` causes the truncation of the high-order 32-bits.

The solution is to store pointers as pointer types or the special types defined for this purpose, such as **`intptr_t`** and **`uintptr_t`**.

Integer “overflow”

- operation results in numeric value that is too large for storage space
- **C99** standard dictates that the result is always modulo, "computation involving unsigned operands can never overflow"
 - **wraparound**: does not overflow into other storage
 - $\text{UINT_MAX} + 1 == 0$




image source: [wikipedia](https://en.wikipedia.org/wiki/Integer_overflow)

Integer “overflow”

- **C99** standard dictates that the result is always modulo, "computation involving unsigned operands can never overflow"
 - what about signed operands?
 - Overflowing a signed integer is an undefined behavior.
 - $\text{INT_MAX} + 1 == ???$

Integer “overflow”

- But for not **C99** settings:
 - **result saturation:**
 - occurs on GPUs and DSP
 - wrap around does not occur, instead a MAXVALUE is always returned
- still does not overflow into adjacent memory




image source: [wikipedia](https://en.wikipedia.org/wiki/Integer_overflow)

Integer “underflow”

- occurs in subtraction
 - unsigned int $x = 0 - 1$
 - $x == 2^{16} - 1$
- **C99** standard dictates that the result is always modulo, for unsigned operands
 - wraparound: does not overflow into other storage
- For signed operands this is undefined

Other Integer Nuances

- `-INT_MIN == Undefined behavior`
- Bit shifting integers:
 - Negative integers cannot be left shifted
 - `-1 << x == Undefined behavior` (for any `x >= 0`)
 - Error to shift a 1 into the sign position
 - `INT_MAX << 1 == Undefined behavior`
 - Error to shift by value `>` than bitwidth of the object
 - x86-32, int is 32 bits. Error to do `(int) x << 33`
 - `x << 32` is ok. equivalent to `x = x xor x`

Other Nuances

- Starting a number with 0 designates octal
 - 1000, 2000, 0100, 0200, 0300, ... 0981
 -

Integer Promotion / Conversion

- unsigned wins

- <https://www.securecoding.cert.org/confluence/display/seccode/INT02-C.+Understand+integer+conversion+rules>
- $(1U > -1) == (1U > UINT_MAX) == 0$

- Operands promoted (up to size) int

- Integer types smaller than int are promoted when an operation is performed on them
 - $(short) x \ll 17$
 - promoted to integer, so this is safe

Other Integer Nuances

- $(\text{int})X - 1 + 1 == \text{undefined}$ IF $X == \text{INT_MIN}$
- $\text{INT_MIN} \% -1$
- Does:
 - $(\text{short})x + 1 == (\text{short})(x+1)$ for all values?

size_t vs ssize_t

size_t = an unsigned int

- rationale: Sizes should never be negative

But stupid things happen:

<http://pubs.opengroup.org/onlinepubs/009604499/functions/mbstowcs.html>

- People want to be able to return (size_t)-1 == -1
 - thus ssize_t

- [-1, SSIZE_MAX]

- SSIZE_MAX = 32 767

*My buddy Sean @ CMU CERT
found this awesome case*

Importance of Integer Bugs

- Crypto Libraries
 - <http://blog.regehr.org/archives/1054>
 - Probably in bitcoin / cryptocurrency libraries
- Often not understood by developers
- Can lead to vulnerabilities
- Suggested reading:
 - <http://www.cs.utah.edu/~regehr/papers/overflow12.pdf>

Floats

Float variable can be NaN (Not a number)

Float Nuances:


- $0.1 + 0.2 = 0.3000000000000000000004$
- `NaN == NaN` is always false
- sums of many floats dont always add up right
 - precision is lost
- Suggested Reading <http://floating-point-gui.>

Bug time! (30 Seconds)

```
// Coefficients USED TO CONVERT FROM RGB TO monochrome.  
const uint32 kRedCoefficient = 2125;  
const uint32 kGreenCoefficient = 7154;  
const uint32 kBlueCoefficient = 0721;  
const uint32 kColorCoefficientDenominator = 10000;
```



This error was found in
theChromium project by
[PVS-Studio](#) C/C++ static
code analyzer.



Integer bug resources

<http://www.ibm.com/developerworks/library/l-port64/>

promotion rules: [https://www.securecoding.cert.org/confluence/display/seccode/INT02-C.](https://www.securecoding.cert.org/confluence/display/seccode/INT02-C)

[+Understand+integer+conversion+rules](#)

floats: <http://floating-point-gui.de/>

crypto libraries: <http://blog.regehr.org/archives/1054>

Formatted Output Security

- Section 0x352 (HAOE) covers this very well
- The problem of Format Strings
 - misuse
 - exploitation
 - Crashing
 - information leak/disclosure
- Mitigation Techniques
 - user input \neq format string

Format Strings

- printf
- sprintf
- snprintf
- fprintf
- syslog
- ...

Looking at how functions are called

Code editor

```
1 // Type your code here, or load an example.
2 #include <stdio.h>
3 void foo(){
4     char buffer[256];
5     sprintf(buffer, "Hello World! #%d", 12345);
6 }
```

Assembly output

```
1 .LC0:
2     .string "Hello World! #%d"
3 foo():
4     push    rbp
5     mov     rbp, rsp
6     sub     rsp, 256
7     lea     rax, [rbp-256]
8     mov     edx, 12345
9     mov     esi, OFFSET FLAT:.LC0
10    mov     rdi, rax
11    mov     eax, 0
12    call    sprintf
13    leave
14    ret
15
```

- arguments passed via registers
 - we'll cover this more next time
- then call sprintf

Looking at how functions are called

64 bit

- using registers

Code editor

```
1 #include <math.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 void foo() {
6     char buffer[256];
7     sprintf(buffer, "Hello %d %d %d %d", 1,2,3,4);
8 }
9
```

Assembly output

```
1 .LC0:
2     .string "Hello %d %d %d %d"
3 foo():
4     push    rbp
5     mov rbp, rsp
6     sub rsp, 256
7     lea rax, [rbp-256]
8     mov r9d, 4
9     mov r8d, 3
10    mov ecx, 2
11    mov edx, 1
12    mov esi, OFFSET FLAT:.LC0
13    mov rdi, rax
14    mov eax, 0
15    call    sprintf
16    leave
17    ret
```

Looking at how functions are called

32 bit (-m32 compiler flag)

- arguments on the stack

Code editor

```
1 #include <math.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 void foo() {
6     char buffer[256];
7     sprintf(buffer, "Hello %d %d %d %d", 1,2,3,4);
8
9 }
```

Assembly output

```
1 .LC0:
2     .string "Hello %d %d %d %d"
3 foo():
4     push    ebp
5     mov     ebp, esp
6     sub     esp, 296
7     mov     DWORD PTR [esp+20], 4
8     mov     DWORD PTR [esp+16], 3
9     mov     DWORD PTR [esp+12], 2
10    mov     DWORD PTR [esp+8], 1
11    mov     DWORD PTR [esp+4], OFFSET FLAT:.LC0
12    lea     eax, [ebp-264]
13    mov     DWORD PTR [esp], eax
14    call    sprintf
15    leave
16    ret
```

Looking at how functions are called

64 bit

- eventually will use the stack

Code editor

```
1 #include <math.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 void foo() {
6     char buffer[256];
7     sprintf(buffer, "Hello %d %d %d %d %d %d %d %d", 1,2,3,4, 5, 6 , 7, 8);
8 }
9
```

Assembly output

```
1 .LC0:
2     .string "Hello %d %d %d %d %d %d %d %d"
3 foo():
4     push    rbp
5     mov     rbp, rsp
6     sub     rsp, 288
7     lea     rax, [rbp-256]
8     mov     DWORD PTR [rsp+24], 8
9     mov     DWORD PTR [rsp+16], 7
10    mov     DWORD PTR [rsp+8], 6
11    mov     DWORD PTR [rsp], 5
12    mov     r9d, 4
13    mov     r8d, 3
14    mov     ecx, 2
15    mov     edx, 1
16    mov     esi, OFFSET FLAT:.LC0
17    mov     rdi, rax
18    mov     eax, 0
19    call    sprintf
20    leave
21    ret
```


Looking at how functions are called

- Depends on architecture
- Depends on calling standard
 - more on this later
- Depends on type of function
 - normal code vs. system call
 - more on this later


Looking at how functions are called

32 bit

- format string function parses these to determine what to use on the stack

Code editor

```
1 #include <math.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 void foo() {
6     char buffer[256];
7     sprintf(buffer, "Hello %d %d %d %d", 1,2,3,4);
8 }
9
```



Assembly output

```
1 .LC0:
2     .string "Hello %d %d %d %d"
3 foo():
4     push    ebp
5     mov     ebp, esp
6     sub     esp, 296
7     mov     DWORD PTR [esp+20], 4
8     mov     DWORD PTR [esp+16], 3
9     mov     DWORD PTR [esp+12], 2
10    mov     DWORD PTR [esp+8], 1
11    mov     DWORD PTR [esp+4], OFFSET FLAT:.LC0
12    lea     eax, [ebp-264]
13    mov     DWORD PTR [esp], eax
14    call    sprintf
15    leave
16    ret
```

Looking at how functions are called

32 bit

- format string ****SAFE**** example

Code editor

```
1 #include <math.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 void foo(char *buffer) {
6     printf("%s", buffer);
7 }
```

Assembly output

```
1 .LC0:
2     .string "%s"
3 foo(char*):
4     push    ebp
5     mov     ebp, esp
6     sub     esp, 24
7     mov     eax, DWORD PTR [ebp+8]
8     mov     DWORD PTR [esp+4], eax
9     mov     DWORD PTR [esp], OFFSET FLAT:.LC0
10    call    printf
11    leave
12    ret
```

Looking at how functions are called

32 bit

- format string vulnerability example

Code editor

```
1 #include <math.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 void foo(char *buffer) {
6     printf(buffer);
7 }
```

Assembly output

```
1 foo(char*):
2     push    ebp
3     mov     ebp, esp
4     sub     esp, 24
5     mov     eax, DWORD PTR [ebp+8]
6     mov     DWORD PTR [esp], eax
7     call    printf
8     leave
9     ret
```

Format Strings

`%[flags][width][.precision][{length-modifier}]` conversion-specifier

- `%d` or `%i` = signed decimal integer
- `%u` = unsigned decimal integer
- `%o` = unsigned octal
- `%x` = unsigned hexadecimal integer
- `%X` = unsigned hexadecimal integer (uppercase)
- `%f` = decimal float
- `%e` = scientific notation
- `%a` = hexadecimal floating point
- `%c` = char
- `%s` = string
- `%p` = pointer address
- `%n` = nothing printed, but corresponds to a pointer. The number of characters written so far is stored in the pointed location.

| | <u>specifiers</u> | | | | | | |
|---------------|-------------------|------------------------|-----------------|--------|----------|-------|----------------|
| <i>length</i> | d i | u o x X | f F e E g G a A | c | s | p | n |
| <i>(none)</i> | int | unsigned int | double | int | char* | void* | int* |
| hh | signed char | unsigned char | | | | | signed char* |
| h | short int | unsigned short int | | | | | short int* |
| l | long int | unsigned long int | | wint_t | wchar_t* | | long int* |
| ll | long long int | unsigned long long int | | | | | long long int* |
| j | intmax_t | uintmax_t | | | | | intmax_t* |
| z | size_t | size_t | | | | | size_t* |
| t | ptrdiff_t | ptrdiff_t | | | | | ptrdiff_t* |
| L | | | long double | | | | |

Exploiting Format Strings

Back to that example:

```
gets(buffer); // buffer == "%s%s..."
```

```
printf(buffer);
```

```
printf("%s%s%s%s%s%s%s%s%s%s...");
```

- reads pointer values off the stack for each %s
 - until all %s specifiers are satisfied
 - or until segfault

Exploiting Format Strings

```
printf("%08x %08x %08x %08x %08x....");
```

- prints out values on the stack in hex format
 - allows viewing of stack contents by attacker
 - printed in human-friendly format
 - x86-64 / x86 values are stored little-endian in memory
 - very important to remember

Exploiting Format Strings

```
printf(“%08x %08x %08x %08x %08x....”);
```

- Iteratively increases the argument pointer by 8 each time.
 - for variable argument functions
 - `va_start`
 - `va_list`
 - an array. `va_list[i]` is argument pointer. (YMMV)

Exploiting Format Strings

```
printf("%04x....");
```

- can move forward argument pointer by other values.
 - typically just by 4 or 8 bytes on x86-32. Not sure on x86-64
 - This can be exploited to view arbitrary memory locations

Exploiting Format Strings

```
printf("\xde\x5\x04%x%x%x%x%s");
```

- viewing arbitrary memory locations (32bit)
 - move argument pointer forward enough to point within the string (the %x chain)
- %s uses a stack value as a pointer
 - prints out what it points to
 - here, will print the value at 04e5f5de (little endian)

Exploiting Format Strings

```
printf("\xde\xef\xef\x04%x%x%x%x%s");
```

- `\x` ← these are escape characters
 - denotes special character
 - ASCII encoding
 - used here to provide a little-endian address (32 bit example)
 - (more on this in the exploitation section)

Exploiting Format Strings

Writing to memory address (from [1] p326)

```
int i;
```

```
printf("hello%n\n", (int *)&i);
```

writes 5 to variable i;

Exploiting Format Strings

Writing to arbitrary memory address

```
printf("\xde\xef\x04%x%x%x%x\n");
```

```
printf("\xde\xef\x04%x%x%x%150x\n");
```

works well for writing small values

- but not memory addresses

Exploiting Format Strings

Writing to arbitrary memory address

```
printf("\xde\xef\x04%x%x%x%x\n");
```

will write the number of characters before the %
n printed so far to 04e5f5de.

- We need to explore length modifier:

%[flags][width][.precision][{length-modifier}] conversion-specifier

| | specifiers | | | | | | |
|---------------|---------------|------------------------|-----------------|--------|----------|-------|----------------|
| <i>length</i> | d i | u o x X | f F e E g G a A | c | s | p | n |
| (none) | int | unsigned int | double | int | char* | void* | int* |
| hh | signed char | unsigned char | | | | | signed char* |
| h | short int | unsigned short int | | | | | short int* |
| l | long int | unsigned long int | | wint_t | wchar_t* | | long int* |
| ll | long long int | unsigned long long int | | | | | long long int* |
| j | intmax_t | uintmax_t | | | | | intmax_t* |
| z | size_t | size_t | | | | | size_t* |
| t | ptrdiff_t | ptrdiff_t | | | | | ptrdiff_t* |
| L | | | long double | | | | |

Other modifiers

`%[flags][width][.precision][{length-modifier}]` conversion-specifier

- width
- precision

```
int i;  
printf("%50u\n", 1, &i); // i = 50  
printf("%5000u\n", 1, &i); // i = 5000
```

Exploiting Format Strings

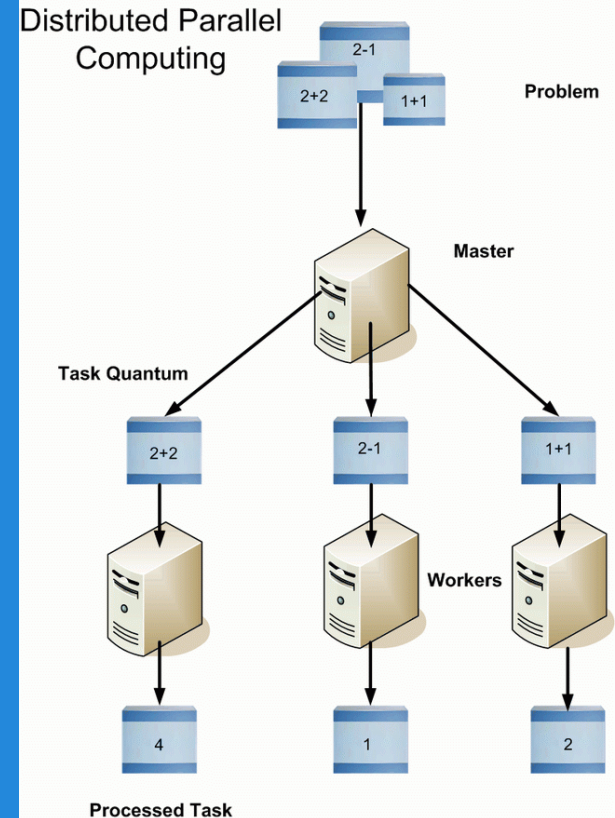
Combine these techniques to write arbitrary values to arbitrary memory location (s) byte by byte:

- pg 174 HAOE explains this best. The following writes 0xDDCCBBAA to the address at 0x08409755

We'll finish
this topic
later in the semester

| <u>Memory</u> | 94 | 95 | 96 | 97 | | | |
|----------------------------|----|----|----|----|----|----|----|
| First write to 0x08409755 | AA | 00 | 00 | 00 | | | |
| Second write to 0x08409756 | | BB | 00 | 00 | 00 | | |
| Third write to 0x08409757 | | | CC | 00 | 00 | 00 | |
| Fourth write to 0x08409758 | | | | DD | 00 | 00 | 00 |
| RESULT | AA | BB | CC | DD | | | |

Concurrency & Race Conditions



Concurrency, Parallelism, & Multithreading

Concurrency:

- Several computations executing simultaneously and potentially interacting with each other
- Concurrency not always equal multithreading
 - *possible for multithreaded applications to not be concurrent*

Multithreading:

- program has two or more threads that **may** execute concurrently

Parallelism:

- Data parallelism vs. task parallelism
 - **data**: split data set into segments
apply function in parallel
 - **task**: split job into several distinct tasks to be run in parallel

3 Properties for Race Conditions [1]

1. Concurrency Property

- At least 2 control flows must be executing concurrently

2. Shared Object Property

- a shared race object must be accessed by both of the concurrent flows

3. Change State Property

- At least one of the control flows must alter the state of the race object

Race Condition explained

concurrency property:

- train
- tracks

shared object

- the tracks (junction)

change state:

- the junction



Race Condition Bughunting Strategy

1. Focus on the shared objects.
2. For each shared object, follow how it is handled through the code. Focus on any state changes.
3. For each state change, enumerate what other concurrent entities might be operating on it.
 - a. Hunt for what could go wrong line by line between the two threads
 - i. R & Ws



Race Condition potential results

- Corrupted Values
- 'Volatile' objects act in undefined ways when handled asynchronously
- Elevated permissions
 - (permission escalation)
 - CVE-2007-4303, CVE-2007-4302, ...
- Deadlocks
 - DoS

Questions?

Reading: 0x280 up to 0x300 (HAOE)
and 0x350 up to 0x400

Great Study Resource: <http://q.viva64.com/>

A diagram consisting of ten blue arrows pointing towards the URL 'http://q.viva64.com/'. The arrows originate from various directions: two from the top, two from the top-right, one from the right, two from the bottom-right, one from the bottom, and two from the bottom-left.